

An AI requirement that can't **fail** can't be validated.

Most AI requirements are aspirations wearing
the word “shall”.

You've probably written this one:

“The system shall accurately extract data from batch records.”

Nothing in that sentence can fail. So nothing in it can be tested — and the tester finds that out during execution, under deadline.

Why it happens:

Nobody knows what number to write. And writing a number feels like committing to something you might miss.

So the requirement stays vague — and the decision quietly moves downstream to whoever is holding the script at 6pm.

THE TEST

Can this requirement fail?

If you can't describe the observable outcome that counts as failure, you don't have a requirement. You have an intention.

Apply it to every line

What a testable AI requirement needs:

- The property — accuracy, recall, consistency, latency
- The threshold — a number, and a reason for it
- The conditions — which data, which population, which document type
- The evidence — how pass and fail get recorded
- The owner — who decides, before anyone runs it

EXAMPLE 01 · EXTRACTION

Before → After

~~The system shall accurately extract data from batch records.~~

≥98% field-level accuracy against a 200-record human-adjudicated reference set. No single field below 95%. Re-measured at qualification and at every model change.

EXAMPLE 02 · HUMAN REVIEW

Before → After

~~AI output shall be reviewed by a qualified user.~~

100% of Tier 3 outputs reviewed by a named role before entering the record. Reviewer disagreement rate monitored — a sustained rate below 2% triggers an automation-bias investigation.

That second one matters more than it looks.

If your reviewers never disagree with the model, they aren't reviewing it.

“Human in the loop” only becomes a control the moment it can fail.

The hardest part is the **threshold**.

Here's the move nobody makes: measure what your current process actually achieves.

If your human reviewers catch 94%, then debating in the abstract whether 97% is “good enough” misses that you already accepted 94% — and never wrote it down.

Setting the number **is** the risk assessment.

To choose a threshold you have to say what happens when it's wrong, who catches it, and what it costs.

That's the conversation everyone says they want and then skips — because it's easier to write “accurately”.

Has your org started the AI-in-GxP journey?

Forward this to your QA, CSV or IT lead — especially if requirements are being drafted right now.

And if a threshold in your own work is hard to justify, ask me.
Happy to think it through.

Sreejith Kanhirangadan

20 yrs CSV · big, mid & startup · 10,000+ trained
evolifeval.com